

# FASHION CLASS CLASSIFICATION

Meet Shah

## 1. PROBLEM STATEMENT AND BUSINESS CASE UNDERSTANDING

Fashion training set consists of 70,000 images divided into 60,000 training and 10,000 testing samples. Dataset sample consists of 28x28 grayscale image, associated with a label from 10 classes.

The 10 classes are as follows:

0 => T-shirt/top 1 => Trouser 2 => Pullover 3 => Dress 4 => Coat 5 => Sandal 6 => Shirt 7 => Sneaker 8 => Bag 9 => Ankle boot

Each image is 28 pixels in height and 28 pixels in width, for a total of 784 pixels in total. Each pixel has a single pixel-value associated with it, indicating the lightness or darkness of that pixel, with higher numbers meaning darker. This pixel-value is an integer between 0 and 255.



## 2. Importing Essential Libraries & Our Fashion Mnist Data

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
#import plotly.express as px
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: train_df = pd.read_csv('fashion-mnist_train.csv', sep = ',')
test_df = pd.read_csv('fashion-mnist_test.csv', sep = ',')
```

```
In [3]: train_df.head(10)
```

```
Out[3]:
```

	label	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	pixel9	...	pixel775	pixel77
0	2	0	0	0	0	0	0	0	0	0	...	0	
1	9	0	0	0	0	0	0	0	0	0	...	0	
2	6	0	0	0	0	0	0	0	5	0	...	0	

	label	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	pixel9	...	pixel775	pixel77
<b>3</b>	0	0	0	0	1	2	0	0	0	0	...	3	
<b>4</b>	3	0	0	0	0	0	0	0	0	0	...	0	
<b>5</b>	4	0	0	0	5	4	5	5	3	5	...	7	
<b>6</b>	4	0	0	0	0	0	0	0	0	0	...	14	
<b>7</b>	5	0	0	0	0	0	0	0	0	0	...	0	
<b>8</b>	4	0	0	0	0	0	0	3	2	0	...	1	
<b>9</b>	8	0	0	0	0	0	0	0	0	0	...	203	21

10 rows × 785 columns

In [4]: `test_df.head(10)`

	label	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	pixel9	...	pixel775	pixel77
<b>0</b>	0	0	0	0	0	0	0	0	9	8	...	103	8
<b>1</b>	1	0	0	0	0	0	0	0	0	0	...	34	
<b>2</b>	2	0	0	0	0	0	0	14	53	99	...	0	
<b>3</b>	2	0	0	0	0	0	0	0	0	0	...	137	12
<b>4</b>	3	0	0	0	0	0	0	0	0	0	...	0	
<b>5</b>	2	0	0	0	0	0	44	105	44	10	...	105	6
<b>6</b>	8	0	0	0	0	0	0	0	0	0	...	0	
<b>7</b>	6	0	0	0	0	0	0	0	1	0	...	174	13
<b>8</b>	5	0	0	0	0	0	0	0	0	0	...	0	
<b>9</b>	0	0	0	0	0	0	0	0	0	0	...	57	7

10 rows × 785 columns

In [5]: `train_df.shape`

Out[5]: (60000, 785)

In [6]: `test_df.shape`

Out[6]: (10000, 785)

## 3. Data Visualization


Converting our data from 'dataframe' to an 'array' so that we can visualize our data.

In [7]: `training = np.array(train_df, dtype = 'float32')`

In [8]: `test = np.array(test_df, dtype = 'float32')`


```
In [9]: plt.imshow(training[55, 1:].reshape(28,28))  
        #plt.savefig('Example Image.png')
```

Out[9]: <matplotlib.image.AxesImage at 0x1b2b7395f40>



Moreover, we can use a for loop and create a grid of random images.

```
In [10]: W_grid = 15  
         L_grid = 15  
         fig, axes = plt.subplots(L_grid, W_grid, figsize = (17,17))  
         axes = axes.ravel()  
         n_training = len(training)  
         for i in np.arange(0, W_grid * L_grid):  
             index = np.random.randint(0, n_training)  
             axes[i].imshow(training[index, 1:].reshape(28,28))  
             axes[i].set_title(training[index, 0], fontsize = 8)  
             axes[i].axis('off')  
         plt.subplots_adjust(hspace=0.4)  
         #fig.savefig('Grid Image.png')
```



## Model Building & Training

```
In [11]: # Splitting the Training Set
X_train = training[:, 1:]/255
y_train = training[:, 0]
```

```
# Splitting the Test Set
X_test = test[:, 1:]/255
y_test = test[:, 0]
```

```
In [12]: from sklearn.model_selection import train_test_split
X_train, X_validate, y_train, y_validate = train_test_split(X_train, y_train, test_s
```

```
In [13]: X_train.shape
```

```
Out[13]: (48000, 784)
```

```
In [14]: y_train.shape
```

```
Out[14]: (48000,)
```

```
In [15]: # Remember that * unpacks the tuple
```

```
X_train = X_train.reshape(X_train.shape[0], *(28, 28 ,1))
X_test = X_test.reshape(X_test.shape[0], *(28, 28 ,1))
X_validate = X_validate.reshape(X_validate.shape[0], *(28, 28 ,1))
```

In [16]: X\_train.shape

Out[16]: (48000, 28, 28, 1)

In [17]: X\_test.shape

Out[17]: (10000, 28, 28, 1)

In [18]: X\_validate.shape

Out[18]: (12000, 28, 28, 1)

In [19]: import tensorflow

In [20]: tensorflow.\_\_version\_\_

Out[20]: '2.3.1'

In [21]: from tensorflow.keras.models import Sequential  
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dense, Flatten, Dropout  
from tensorflow.keras.optimizers import Adam  
from tensorflow.keras.callbacks import TensorBoard

In [22]: cnn = Sequential()

In [23]: *# Was first run with filter = 32 and then later changed it to 64 for better accuracy*  
cnn.add(Conv2D(filters = 64, kernel\_size = 3, input\_shape = [28, 28, 1], activation

In [24]: cnn.add(MaxPooling2D(pool\_size = (2, 2)))

In [25]: cnn.add(Flatten())

In [26]: cnn.add(Dense(units = 32, activation = 'relu'))

In [27]: cnn.add(Dense(units = 10, activation = 'sigmoid'))

In [ ]: *#cnn.add(Dropout(0.25))*

In [28]: cnn.compile(loss = 'sparse\_categorical\_crossentropy', optimizer = Adam(lr=0.001), me

In [29]: cnn.fit(X\_train, y\_train, epochs = 50, batch\_size = 512, verbose = 1, validation\_dat

```
Epoch 1/50
94/94 [=====] - 41s 431ms/step - loss: 0.8849 - accuracy:
0.6748 - val_loss: 0.4992 - val_accuracy: 0.8219
Epoch 2/50
94/94 [=====] - 34s 364ms/step - loss: 0.4484 - accuracy:
0.8426 - val_loss: 0.4152 - val_accuracy: 0.8547
Epoch 3/50
94/94 [=====] - 34s 361ms/step - loss: 0.3906 - accuracy:
0.8638 - val_loss: 0.3742 - val_accuracy: 0.8714
Epoch 4/50
94/94 [=====] - 35s 376ms/step - loss: 0.3556 - accuracy:
0.8764 - val_loss: 0.3692 - val_accuracy: 0.8717
Epoch 5/50
```

94/94 [=====] - 34s 361ms/step - loss: 0.3331 - accuracy: 0.8849 - val\_loss: 0.3280 - val\_accuracy: 0.8832  
Epoch 6/50  
94/94 [=====] - 34s 363ms/step - loss: 0.3139 - accuracy: 0.8901 - val\_loss: 0.3213 - val\_accuracy: 0.8885  
Epoch 7/50  
94/94 [=====] - 35s 371ms/step - loss: 0.3022 - accuracy: 0.8943 - val\_loss: 0.3119 - val\_accuracy: 0.8907  
Epoch 8/50  
94/94 [=====] - 34s 358ms/step - loss: 0.2880 - accuracy: 0.8996 - val\_loss: 0.3050 - val\_accuracy: 0.8939  
Epoch 9/50  
94/94 [=====] - 34s 366ms/step - loss: 0.2782 - accuracy: 0.9022 - val\_loss: 0.2902 - val\_accuracy: 0.8979  
Epoch 10/50  
94/94 [=====] - 34s 363ms/step - loss: 0.2670 - accuracy: 0.9062 - val\_loss: 0.2968 - val\_accuracy: 0.8951  
Epoch 11/50  
94/94 [=====] - 34s 364ms/step - loss: 0.2608 - accuracy: 0.9079 - val\_loss: 0.2808 - val\_accuracy: 0.9022  
Epoch 12/50  
94/94 [=====] - 34s 364ms/step - loss: 0.2502 - accuracy: 0.9119 - val\_loss: 0.2795 - val\_accuracy: 0.9021  
Epoch 13/50  
94/94 [=====] - 33s 348ms/step - loss: 0.2441 - accuracy: 0.9138 - val\_loss: 0.2921 - val\_accuracy: 0.8962  
Epoch 14/50  
94/94 [=====] - 33s 356ms/step - loss: 0.2398 - accuracy: 0.9141 - val\_loss: 0.2721 - val\_accuracy: 0.9038  
Epoch 15/50  
94/94 [=====] - 34s 356ms/step - loss: 0.2289 - accuracy: 0.9197 - val\_loss: 0.2735 - val\_accuracy: 0.9023  
Epoch 16/50  
94/94 [=====] - 36s 385ms/step - loss: 0.2252 - accuracy: 0.9210 - val\_loss: 0.2593 - val\_accuracy: 0.9084  
Epoch 17/50  
94/94 [=====] - 34s 362ms/step - loss: 0.2164 - accuracy: 0.9243 - val\_loss: 0.2598 - val\_accuracy: 0.9078  
Epoch 18/50  
94/94 [=====] - 33s 356ms/step - loss: 0.2134 - accuracy: 0.9249 - val\_loss: 0.2567 - val\_accuracy: 0.9087  
Epoch 19/50  
94/94 [=====] - 34s 357ms/step - loss: 0.2111 - accuracy: 0.9255 - val\_loss: 0.2551 - val\_accuracy: 0.9090  
Epoch 20/50  
94/94 [=====] - 33s 356ms/step - loss: 0.2013 - accuracy: 0.9291 - val\_loss: 0.2563 - val\_accuracy: 0.9078  
Epoch 21/50  
94/94 [=====] - 33s 355ms/step - loss: 0.2027 - accuracy: 0.9281 - val\_loss: 0.2604 - val\_accuracy: 0.9081  
Epoch 22/50  
94/94 [=====] - 33s 355ms/step - loss: 0.1920 - accuracy: 0.9321 - val\_loss: 0.2521 - val\_accuracy: 0.9093  
Epoch 23/50  
94/94 [=====] - 34s 365ms/step - loss: 0.1875 - accuracy: 0.9339 - val\_loss: 0.2531 - val\_accuracy: 0.9084  
Epoch 24/50  
94/94 [=====] - 33s 356ms/step - loss: 0.1832 - accuracy: 0.9364 - val\_loss: 0.2563 - val\_accuracy: 0.9094  
Epoch 25/50  
94/94 [=====] - 33s 356ms/step - loss: 0.1789 - accuracy: 0.9383 - val\_loss: 0.2498 - val\_accuracy: 0.9115  
Epoch 26/50  
94/94 [=====] - 33s 355ms/step - loss: 0.1757 - accuracy: 0.9381 - val\_loss: 0.2470 - val\_accuracy: 0.9120  
Epoch 27/50  
94/94 [=====] - 33s 355ms/step - loss: 0.1675 - accuracy: 0.9421 - val\_loss: 0.2486 - val\_accuracy: 0.9128  
Epoch 28/50

94/94 [=====] - 33s 356ms/step - loss: 0.1674 - accuracy:  
0.9408 - val\_loss: 0.2552 - val\_accuracy: 0.9125  
Epoch 29/50  
94/94 [=====] - 33s 355ms/step - loss: 0.1638 - accuracy:  
0.9422 - val\_loss: 0.2497 - val\_accuracy: 0.9114  
Epoch 30/50  
94/94 [=====] - 33s 356ms/step - loss: 0.1580 - accuracy:  
0.9440 - val\_loss: 0.2635 - val\_accuracy: 0.9093  
Epoch 31/50  
94/94 [=====] - 33s 355ms/step - loss: 0.1543 - accuracy:  
0.9458 - val\_loss: 0.2558 - val\_accuracy: 0.9081  
Epoch 32/50  
94/94 [=====] - 34s 365ms/step - loss: 0.1506 - accuracy:  
0.9477 - val\_loss: 0.2515 - val\_accuracy: 0.9103  
Epoch 33/50  
94/94 [=====] - 33s 349ms/step - loss: 0.1479 - accuracy:  
0.9483 - val\_loss: 0.2482 - val\_accuracy: 0.9139  
Epoch 34/50  
94/94 [=====] - 33s 354ms/step - loss: 0.1466 - accuracy:  
0.9490 - val\_loss: 0.2708 - val\_accuracy: 0.9063  
Epoch 35/50  
94/94 [=====] - 33s 355ms/step - loss: 0.1404 - accuracy:  
0.9508 - val\_loss: 0.2637 - val\_accuracy: 0.9089  
Epoch 36/50  
94/94 [=====] - 33s 352ms/step - loss: 0.1365 - accuracy:  
0.9532 - val\_loss: 0.2634 - val\_accuracy: 0.9093  
Epoch 37/50  
94/94 [=====] - 33s 347ms/step - loss: 0.1343 - accuracy:  
0.9532 - val\_loss: 0.2572 - val\_accuracy: 0.9103  
Epoch 38/50  
94/94 [=====] - 34s 364ms/step - loss: 0.1294 - accuracy:  
0.9551 - val\_loss: 0.2561 - val\_accuracy: 0.9128  
Epoch 39/50  
94/94 [=====] - 33s 355ms/step - loss: 0.1292 - accuracy:  
0.9543 - val\_loss: 0.2584 - val\_accuracy: 0.9133  
Epoch 40/50  
94/94 [=====] - 34s 358ms/step - loss: 0.1240 - accuracy:  
0.9568 - val\_loss: 0.2574 - val\_accuracy: 0.9138  
Epoch 41/50  
94/94 [=====] - 34s 359ms/step - loss: 0.1203 - accuracy:  
0.9590 - val\_loss: 0.2619 - val\_accuracy: 0.9113  
Epoch 42/50  
94/94 [=====] - 34s 362ms/step - loss: 0.1166 - accuracy:  
0.9602 - val\_loss: 0.2631 - val\_accuracy: 0.9118  
Epoch 43/50  
94/94 [=====] - 34s 360ms/step - loss: 0.1146 - accuracy:  
0.9604 - val\_loss: 0.2787 - val\_accuracy: 0.9084  
Epoch 44/50  
94/94 [=====] - 33s 352ms/step - loss: 0.1123 - accuracy:  
0.9617 - val\_loss: 0.2659 - val\_accuracy: 0.9139  
Epoch 45/50  
94/94 [=====] - 36s 381ms/step - loss: 0.1059 - accuracy:  
0.9648 - val\_loss: 0.2669 - val\_accuracy: 0.9127  
Epoch 46/50  
94/94 [=====] - 34s 362ms/step - loss: 0.1077 - accuracy:  
0.9640 - val\_loss: 0.2693 - val\_accuracy: 0.9121  
Epoch 47/50  
94/94 [=====] - 33s 351ms/step - loss: 0.1034 - accuracy:  
0.9648 - val\_loss: 0.2665 - val\_accuracy: 0.9135  
Epoch 48/50  
94/94 [=====] - 33s 356ms/step - loss: 0.1005 - accuracy:  
0.9659 - val\_loss: 0.3021 - val\_accuracy: 0.9014  
Epoch 49/50  
94/94 [=====] - 33s 352ms/step - loss: 0.1001 - accuracy:  
0.9656 - val\_loss: 0.2731 - val\_accuracy: 0.9146  
Epoch 50/50  
94/94 [=====] - 33s 352ms/step - loss: 0.0950 - accuracy:  
0.9679 - val\_loss: 0.2716 - val\_accuracy: 0.9162

Out[29]: <tensorflow.python.keras.callbacks.History at 0x1b2f6313ee0>



```
In [30]: evaluation = cnn.evaluate(X_test, y_test)
print('Test Accuracy : {:.3f}'.format(evaluation[1]))

313/313 [=====] - 3s 8ms/step - loss: 0.2745 - accuracy: 0.9172
Test Accuracy : 0.917
```


```
In [31]: predicted_classes = cnn.predict_classes(X_test)

WARNING:tensorflow:From <ipython-input-31-ef5bdc73df9d>:1: Sequential.predict_classes (from tensorflow.python.keras.engine.sequential) is deprecated and will be removed after 2021-01-01.
Instructions for updating:
Please use instead: * `np.argmax(model.predict(x), axis=-1)`, if your model does multi-class classification (e.g. if it uses a `softmax` last-layer activation). * `(model.predict(x) > 0.5).astype("int32")`, if your model does binary classification (e.g. if it uses a `sigmoid` last-layer activation).
```


```
In [32]: predicted_classes
```

```
Out[32]: array([0, 1, 2, ..., 8, 8, 1], dtype=int64)
```

```
In [33]: L = 5
W = 5
fig, axes = plt.subplots(L, W, figsize = (12,12))
axes = axes.ravel()
for i in np.arange(0, L * W):
    axes[i].imshow(X_test[i].reshape(28, 28))
    axes[i].set_title("Prediction Class = {:.1f}\n True Class = {:.1f}".format(predicted_classes[i], y_test[i]))
    axes[i].axis('off')
plt.subplots_adjust(wspace = 0.5)
#fig.savefig('Grid Image Showing Predicted Class vs True Class.png')
```



```
In [34]: from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, predicted_classes)
plt.figure(figsize = (14, 10))
sns.heatmap(cm, annot = True)
#plt.savefig('Heatmap Showing Predicted class vs True Class.png')
```



```
In [35]: from sklearn.metrics import classification_report

target_names = ["Class {}".format(i) for i in range(1, 11)]
print(classification_report(y_test, predicted_classes, target_names = target_names))
```

	precision	recall	f1-score	support
Class 1	0.85	0.87	0.86	1000
Class 2	0.99	0.99	0.99	1000
Class 3	0.91	0.84	0.87	1000
Class 4	0.92	0.92	0.92	1000
Class 5	0.86	0.91	0.88	1000
Class 6	0.99	0.96	0.98	1000
Class 7	0.77	0.78	0.77	1000
Class 8	0.95	0.96	0.96	1000
Class 9	0.98	0.98	0.98	1000
Class 10	0.96	0.97	0.97	1000
accuracy			0.92	10000
macro avg	0.92	0.92	0.92	10000
weighted avg	0.92	0.92	0.92	10000

## Conclusion

As seen in output no. 29 & 30, our accuracies are as follows:

1. Training Data Accuracy - 96.79%
2. Validation Data Accuracy - 91.62%
3. Testing Data Accuracy - 91.7%

So when we print out our classification report we finally get an average accuracy of 92%.

